



Ask the Experts AMD Playbooks

Benjamin Consolvo, SMTS AI Engineer @ AMD
Sreeram Sivasubramony, ML Solutions Engineer @ AMD
August 13, 2025 | Hosted on AMD Developer Community Discord


Agenda

-
1. What are AMD Playbooks?
 2. 22 Playbooks and counting
 3. Supported Hardware & OSes
 4. Live Demo: Website & Halo Box
 5. Quality: Never Stale with Nightly CI
 6. Live Demo: Generating Images with ComfyUI and Z Image Turbo
 7. Contribute a Playbook
 8. Resources + Q&A

1. What are AMD Playbooks?

Your guided on-ramp to AI on AMD hardware

- Step-by-step guided walkthroughs for running real AI workloads on AMD hardware
- Not a manual, not documentation, not every-single-button — a curated path that gets you productive fast
- Educational: learn something new about an AI model or technique
- Discover the popular open-source frameworks and AI use cases
- Kickstart your developer journey on AMD — pick your device & OS, follow along, and you're running



2. 22 Playbooks and counting

Windows & Linux — laptop APU to desktop GPU

Image, Vision & Agents

- Generating Images with ComfyUI and Z Image Turbo ★
- Local Computer Vision with AMD Ryzen™ AI NPU
- Real-Time Speech-to-Speech Translation
- Building Your First Agent with GAIA
- Running Hermes Agent Locally with Lemonade Server
- Running OpenClaw Locally with Lemonade Server
- Chatting with LLMs in Open WebUI

LLMs, Serving & Dev Tools

- Getting Started with Lemonade
- Getting Started with Ollama
- Getting Started with vLLM
- Running and Serving LLMs with LM Studio
- Running LLMs on PyTorch with AMD ROCm™
- Local LLM Coding with VSCode and Qwen3-Coder
- Remote Development with AMD Sync
- Running DeepSeek V4 Flash with ds4

Fine-Tuning, Kernels & Scaling

- Fine-Tuning LLMs with Unsloth
- Fine-Tuning LLMs with LLaMA Factory
- Fine-Tuning LLMs with PyTorch and ROCm™
- Building Custom GPU Kernels with PyTorch
- Automating Workflows with n8n and Local LLMs
- Clustering Two Ryzen™ AI Halos with RCCL
- Clustering Two Ryzen™ AI Halos with RPC

3. Supported Hardware & OSes

New playbooks added regularly

Hardware families

AMD Ryzen™ AI Halo

- Ryzen AI Max+ — Halo Box: pre-installed experience

AMD Ryzen™ AI APUs

- Ryzen AI Max+ · Ryzen AI 300 HX · Ryzen AI 300

AMD Radeon™ GPUs

- RX 9000 Series · RX 7000 / PRO W7000 · AI Pro R9700

Operating systems: Windows & Linux

Good to know

- Not every playbook runs on every device
- Each playbook lists its supported devices — use the filters on the website (processor, platform, difficulty)
- Example: ComfyUI + Z Image Turbo runs on Ryzen AI Max+ / Halo and Radeon RX 7900 XT, RX 9070 XT, R9700 — but not on Ryzen AI 300
- Halo Box = zero setup: apps pre-installed, models pre-downloaded, one-click launcher



4. Live Demo: Website & Halo Box

5. Quality: Never Stale with Nightly CI

Quality: automated nightly regression on real AMD hardware

- **Every playbook is automatically tested every night** — we run the actual commands from each guide, exactly as you would
- **On real AMD hardware** — Ryzen™ AI and Radeon™ devices, on both Windows and Linux
- **Against the latest framework versions** — pulled fresh each night, so new releases can't silently break a step
- **If a step breaks, we catch and fix it** — before you ever hit it

What this means for you

Open any playbook right now — the steps run **as if written this morning**, not a month ago.



6. Live Demo: Generating Images with ComfyUI and Z Image Turbo

7. Contribute a Playbook

Bring your own ideas to AMD Playbooks!

1. You submit an abstract (~ 10 minutes)
2. AMD playbook maintainers review it and label the issue `status::approved`, `status::needs-info`, or `status::declined`.
3. If it is approved, you open a pull request following the Playbook Creation Guide and link it back to the issue. Pull requests without an approved proposal will be closed.

Find out more at the CONTRIBUTING.md guide:

<https://github.com/amd/playbooks/blob/main/CONTRIBUTING.md>

Create new issue

Templates and forms

Bug report →

Report a problem with a specific playbook in this repository.

Playbook proposal →

Pitch a new playbook as a short abstract, before writing any of it.

8. Resources + Q&A

Visit the links below to get more involved

AMD Playbooks

<https://developer.amd.com/playbooks>



AMD Playbooks GitHub

<https://github.com/amd/playbooks>



AMD AI Developer Program

<https://developer.amd.com/ai-developer-program/>

